

This notebook is an exercise in the [Intro to Deep Learning](#) course. You can reference the tutorial at [this link](#).

Introduction

In the tutorial we learned about the building blocks of neural networks: *linear units*. We saw that a model of just one linear unit will fit a linear function to a dataset (equivalent to linear regression). In this exercise, you'll build a linear model and get some practice working with models in Keras.

Before you get started, run the code cell below to set everything up.

In []:

```
# Setup plotting
import matplotlib.pyplot as plt

plt.style.use('seaborn-whitegrid')
# Set Matplotlib defaults
plt.rc('figure', autolayout=True)
plt.rc('axes', labelweight='bold', labelsize='large',
       titleweight='bold', titlesize=18, titlepad=10)

# Setup feedback system
from learntools.core import binder
binder.bind(globals())
from learntools.deep_learning_intro.ex1 import *
```

The *Red Wine Quality* dataset consists of physiochemical measurements from about 1600 Portuguese red wines. Also included is a quality rating for each wine from blind taste-tests.

First, run the next cell to display the first few rows of this dataset.

In []:

```
import pandas as pd

red_wine = pd.read_csv('../input/dl-course-data/red-wine.csv')
red_wine.head()
```

You can get the number of rows and columns of a dataframe (or a Numpy array) with the `shape` attribute.

In []:

```
red_wine.shape # (rows, columns)
```

1) Input shape

How well can we predict a wine's perceived quality from the physiochemical measurements?

The target is `'quality'`, and the remaining columns are the features. How would you set the `input_shape` parameter for a Keras model on this task?

In []:

```
# YOUR CODE HERE
input_shape = [red_wine.shape[1]-1]

# Check your answer
q_1.check()
```

```
In [ ]:
```

```
# Lines below will give you a hint or solution code
#q_1.hint()
#q_1.solution()
```

2) Define a linear model

Now define a linear model appropriate for this task. Pay attention to how many inputs and outputs the model should have.

```
In [ ]:
```

```
from tensorflow import keras
from tensorflow.keras import layers

# YOUR CODE HERE
# unit -> nb of output
# input_shape -> nb of output
model = keras.Sequential([layers.Dense(units=1, input_shape=input_shape)])
print(model)
# Check your answer
q_2.check()
```

```
In [ ]:
```

```
# Lines below will give you a hint or solution code
#q_2.hint()
#q_2.solution()
```

3) Look at the weights

Internally, Keras represents the weights of a neural network with **tensors** (a multidimensional array). Tensors are basically TensorFlow's version of a Numpy array with a few differences that make them better suited to deep learning. One of the most important is that tensors are compatible with [GPU](#) and [TPU](#) accelerators. TPUs, in fact, are designed specifically for tensor computations.

A model's weights are kept in its `weights` attribute as a list of tensors. Get the weights of the model you defined above. (If you want, you could display the weights with something like:

```
print("Weights\n{}\n\nBias\n{}".format(w, b))
```

```
In [ ]:
```

```
# YOUR CODE HERE
# be careful .get_weights() return a numpy array whereas .weights() return tensors
w, b = model.weights
print("Weights\n{}\n\nBias\n{}".format(w, b))
# Check your answer
q_3.check()
```

```
In [ ]:
```

```
# Lines below will give you a hint or solution code
#q_3.hint()
#q_3.solution()
```

(By the way, Keras represents weights as tensors, but also uses tensors to represent data. When you set the `input_shape` argument, you are telling Keras the dimensions of the array it should expect for each example in the training data. Setting `input_shape=[3]` would create a network accepting vectors of length 3, like `[0.2, 0.4, 0.6]`.)

Optional: Plot the output of an untrained linear model

The kinds of problems we'll work on through Lesson 5 will be *regression* problems where the goal is to predict

The kinds of problems we'll work on through Lesson 2 will be regression problems, where the goal is to predict some numeric target. Regression problems are like "curve-fitting" problems: we're trying to find a curve that best fits the data. Let's take a look at the "curve" produced by a linear model. (You've probably guessed that it's a line!)

We mentioned that before training a model's weights are set randomly. Run the cell below a few times to see the different lines produced with a random initialization. (There's no coding for this exercise -- it's just a demonstration.)

In []:

```
import tensorflow as tf
import matplotlib.pyplot as plt

model = keras.Sequential([
    layers.Dense(1, input_shape=[1]),
])

x = tf.linspace(-1.0, 1.0, 100)
y = model.predict(x)

plt.figure(dpi=100)
plt.plot(x, y, 'k')
plt.xlim(-1, 1)
plt.ylim(-1, 1)
plt.xlabel("Input: x")
plt.ylabel("Target y")
w, b = model.weights # you could also use model.get_weights() here
plt.title("Weight: {:0.2f}\nBias: {:0.2f}".format(w[0][0], b[0]))
plt.show()
```

Keep Going

Add hidden layers and [make your models deep](#) in Lesson 2.

Have questions or comments? Visit the [Learn Discussion forum](#) to chat with other Learners.