

This notebook is an exercise in the [Intro to Deep Learning](#) course. You can reference the tutorial at [this link](#).

Introduction

In the tutorial, we saw how to build deep neural networks by stacking layers inside a `Sequential` model. By adding an *activation function* after the hidden layers, we gave the network the ability to learn more complex (non-linear) relationships in the data.

In these exercises, you'll build a neural network with several hidden layers and then explore some activation functions beyond ReLU. Run this next cell to set everything up!

In []:

```
import tensorflow as tf

# Setup plotting
import matplotlib.pyplot as plt

plt.style.use('seaborn-whitegrid')
# Set Matplotlib defaults
plt.rc('figure', autolayout=True)
plt.rc('axes', labelweight='bold', labelsize='large',
       titleweight='bold', titlesize=18, titlepad=10)

# Setup feedback system
from learntools.core import binder
binder.bind(globals())
from learntools.deep_learning_intro.ex2 import *
```

In the *Concrete* dataset, your task is to predict the compressive strength of concrete manufactured according to various recipes.

Run the next code cell without changes to load the dataset.

In []:

```
import pandas as pd

concrete = pd.read_csv('../input/dl-course-data/concrete.csv')
concrete.head()
```

1) Input Shape

The target for this task is the column `'CompressiveStrength'`. The remaining columns are the features we'll use as inputs.

What would be the input shape for this dataset?

In []:

```
# YOUR CODE HERE
input_shape = [concrete.shape[1]-1]

# Check your answer
q_1.check()
```

In []:

```
# Lines below will give you a hint or solution code
#q_1.hint()
```

```
#q_1.solution()
```

2) Define a Model with Hidden Layers

Now create a model with three hidden layers, each having 512 units and the ReLU activation. Be sure to include an output layer of one unit and no activation, and also `input_shape` as an argument to the first layer.

In []:

```
from tensorflow import keras
from tensorflow.keras import layers

# YOUR CODE HERE
# units = nb of output hidden
model = keras.Sequential([
    layers.Dense(units=512, activation="relu", input_shape=input_shape),
    layers.Dense(units=512, activation="relu"),
    layers.Dense(units=512, activation="relu"),
    layers.Dense(units=1),
])

# Check your answer
q_2.check()
```

In []:

```
# Lines below will give you a hint or solution code
#q_2.hint()
#q_2.solution()
```

3) Activation Layers

Let's explore activations functions some.

The usual way of attaching an activation function to a `Dense` layer is to include it as part of the definition with the `activation` argument. Sometimes though you'll want to put some other layer between the `Dense` layer and its activation function. (We'll see an example of this in Lesson 5 with *batch normalization*.) In this case, we can define the activation in its own `Activation` layer, like so:

```
layers.Dense(units=8),
layers.Activation('relu')
```

This is completely equivalent to the ordinary way: `layers.Dense(units=8, activation='relu')`.

Rewrite the following model so that each activation is in its own `Activation` layer.

In []:

```
### YOUR CODE HERE: rewrite this to use activation layers
model = keras.Sequential([

    layers.Dense(32, input_shape=[8]),
    layers.Activation( activation='relu'),

    layers.Dense(32),
    layers.Activation( activation='relu'),

    layers.Dense(1),
])

# Check your answer
q_3.check()
```

In []:

```
# Lines below will give you a hint or solution code
```

```
#q_3.hint()
#q_3.solution()
```

Optional: Alternatives to ReLU

There is a whole family of variants of the `'relu'` activation -- `'elu'`, `'selu'`, and `'swish'`, among others -- all of which you can use in Keras. Sometimes one activation will perform better than another on a given task, so you could consider experimenting with activations as you develop a model. The ReLU activation tends to do well on most problems, so it's a good one to start with.

Let's look at the graphs of some of these. Change the activation from `'relu'` to one of the others named above. Then run the cell to see the graph. (Check out the [documentation](#) for more ideas.)

In []:

```
# YOUR CODE HERE: Change 'relu' to 'elu', 'selu', 'swish'... or something else
activation_layer = layers.Activation('swish')

x = tf.linspace(-3.0, 3.0, 100)
y = activation_layer(x) # once created, a layer is callable just like a function

plt.figure(dpi=100)
plt.plot(x, y)
plt.xlim(-3, 3)
plt.xlabel("Input")
plt.ylabel("Output")
plt.show()
```

Keep Going

Now move on to Lesson 3 and [learn how to train neural networks](#) with stochastic gradient descent.

Have questions or comments? Visit the [Learn Discussion forum](#) to chat with other Learners.